

- A learning agent can be divided into four conceptual components, as shown in Figure (previous slide).
- The most important distinction is between the **learning element**, which is responsible for making improvements, and the **performance element**, which is responsible for selecting external actions. **The performance element is what we have previously considered to be the entire agent: it takes in percepts and decides on actions.**
- **The learning element uses feedback from the critic on how the agent is doing and determines how the performance element should be modified to do better in the future.**
- The design of the learning element depends very much on the design of the performance element. When trying to design an agent that learns a certain capability, the first question is not “How am I going to get it to learn this?” but “What kind of performance element will my agent need to do this once it has learned how?” **Given an agent design, learning mechanisms can be constructed to improve every part of the agent.**
- **The critic tells the learning element how well the agent is doing with respect to a fixed performance standard.** The critic is necessary because the percepts themselves provide no indication of the agent’s success. For example, a chess program could receive a percept indicating that it has checkmated its opponent, but it needs a performance standard to know that this is a good thing; the percept itself does not say so.
- **It is important that the performance standard be fixed.** Conceptually, one should think of it as being outside the agent altogether because the agent must not modify it to fit its own behavior.
- The last component of the learning agent is the **problem generator**. It is responsible for suggesting actions that will lead to new and informative experiences. The point is that if the performance element had its way, it would keep doing the actions that are best, given what it knows. **But if the agent is willing to explore a little and do some perhaps suboptimal actions in the short run, it might discover much better actions for the long run. The problem generator’s job is to suggest these exploratory actions.**